

Getting Started with Julia: Notes

Moritz M. Konarski

August 25, 2020

Contents

The Rationale for Julia	1
The scope of Julia	1
Julia's place among the other programming languages	1
A comparison with other languages for the data scientist	2
MATLAB	2
R	2
Python	2
Useful Links	2
Installing the Julia Platform	3
Working with Julia's shell	3
Startup option and Julia scripts	4
Packages	4
How Julia works	4
Variables, Types, and Operations	6
Variables, naming conventions, and comments	6
Types	6
Integers and Booleans	7
Floating point numbers	7
Elementary mathematical functions and operations	7
Rational and complex numbers	8
Characters	8
Strings	8
Formatting numbers and strings	9
Regular Expressions	9
Ranges and arrays	10
Other ways to create arrays	11
Some common functions for arrays	11
Array of chars to string	11
Dates and times	12
Scope and constants	12

The Rationale for Julia

The scope of Julia

- born out of frustration with existing tools for technical computing
- prototyping needs a easy-to-use language, flexible, high-level language so the focus may be on the problem
- actual computation needs maximum performance hence for production things tend to be re-written in C or Fortran
- this lead to prototyping in slow but easy languages and then re-writes in difficult but fast languages
- Julia was designed to bridge this gap – using LLVM JIT (Just in Time) makes near-C speeds possible while keeping high-level usability

This resulted in:

- open source and liberal license (MIT)
- easy to use and learn, elegant, clear, dynamic, familiar, almost like pseudocode:

```
x -> 7x^3 + 30x^2 + 5x + 42
```

- Julia provides the needed speed without the need to switch languages
- metaprogramming to increase capability
- useable for normal computing tasks, not just pure computing
- easy-to-use multicore and parallel capabilities

Julia's place among the other programming languages

- Julia brings together the two worlds of typed and untyped languages
- Julia does not have a static compilation step but uses a type-inference engine to nonetheless deliver similar speeds
- types can still be used to make compilation easier and to document the code
- *dynamic multiple dispatch* is the approach to pick the best fitting function out of a pool of functions depending on the data type, it's basically polymorphism with type inference
- Julia does not have static type checking however, runtime errors can occur if data types do not match
- Julia also makes it easy to design pure functions and apply functional programming
- Julia is also suited for general purpose programming similar to Python

A comparison with other languages for the data scientist

- Julia's speed approaches C and leaves all other normal alternatives behind
- one of Julia's goals is that one never has to step down to C
- Julia is especially good at running MATLAB and R style programs

MATLAB

- the syntax should be very familiar for MATLAB users, but Julia is more general purpose
- most function names are similar to MATLAB and not R
- Julia is much faster than MATLAB, but it can also interface with it

R

- until now R has dominated statistics
- Julia has the same level of usability, but is 10 to 1000 times faster
- Julia also has an interface to R

Python

- Julia is again much faster than Python, reads similar to it, and can interface with it

Useful Links

- main website: <http://julialang.org>
- documentation: <http://docs.julialang.org>
- packages: <http://pkg.julialang.org>

Installing the Julia Platform

- if parallelization (n concurrent processes) is to be used, compile the julia code with

```
make -j n
```

Working with Julia's shell

- use `quit()` or `CTRL + D` to quit the REPL
- after an expression is evaluated, the result will be stored in the variable `ans`, but only in REPL
- assign values like so:

```
a = 3
```

- type annotations are not needed, they are inferred
- strings are defined by `"` (double quotes)
- to clear the screen but keep the data or variables, type `CTRL + L`
- to clear the workspace and variables, use `workspace()`
- all previous commands are stored in a `.julia_history` file at `/home/$USER/`
- typing `?` will give access to the docs, specific help is available through `help(<item>)`
- to find all the places a function is defined or used, type `apropos("<name>")`
- multiple commands on one line are separated by `;`
- multi-line expressions also work and the shell will wait until the expression is complete

```
julia> if 10 > 0
        println("10 is bigger than 0")
    end
10 is bigger than 0
julia>
```

- use `tab` for automatic completion, double `tab` to show the available functions
- starting a line with `;` makes the rest of the line a shell command
- to exit shell mode, type `backspace`
- the REPL can also execute written programs with

```
julia> include("<name>.jl")
```

- the content of the file will then be executed
- for keybindings see [here](#)

Startup option and Julia scripts

- commands can be evaluated from the command line without starting the repl

```
julia -e 'a = 6 * 7; println(a)'
```

- a script taking arguments can be run like this

```
julia script.jl arg1 arg2 arg3
```

- the arguments are then available in the global constant `ARGS`
- files can also execute other files by calling `include("file.jl")` in them

Packages

- official Julia packages can be found at METADATA.jl at <https://github.com/JuliaLang/METADATA.jl>
- a searchable list can be found at <http://pkg.julialang.org/>
- Julia has a built-in package manager called `Pkg` for installing packages
- to find out which packages are installed, use `Pkg.status()`
- one of the better packages is IJulia, a jupyter mod

```
using PyPlot
x = linspace(0, 5)
y = cos(2x + 5)
plot(x, y, linewidth=2.0, linestyle="--")
title("a nice cosine")
xlabel("x axis")
ylabel("y axis")
```

How Julia works

- Julia uses LLVM JIT to generate machine code just-in-time
- the process works like this:
 1. when a function is run the types are inferred
 2. the JIT compiler turns the function into native machine code
 3. the next time a function is called the already compiled code is run (this is the reason that functions are faster the second time around – important for benchmarking)
- the code is dynamic because it is not dependent on the type of the variable
- these functions are by default *generic*, but JIT bytecode for specific types can be inspected like this

```
julia> f(x) = 2x + 5
      f(generic function with 1 method)
julia> code_llvm(f, (Int64,))
```

- the same can be done to inspect the assembly code using the function `code_native(f, (Int64,))`
- Julia automatically allocates and frees memory, it has a GC that runs at the same time as the program and is somewhat unpredictable

- calling `gc()` will call the GC, `gc_disable()` to disable it

Variables, Types, and Operations

- Julia is an optionally typed language – users can choose to specify the types
- typing in Julia is important for speed, documentation, tooling

Variables, naming conventions, and comments

- Julia differentiates between strings and characters, strings are denoted by double quotes, while characters are denoted by single quotes
- to see what type a variable or reference is one can use the `typeof(<var>)` function
- variables don't have to be typed, but they have to be initialized
- variables can change type, they can be over-written
- **everything** is an expression in Julia
- Julia is strongly typed
- variable names have to begin with a letter, then it can be letter, number, underscore, exclamation point, including unicode characters
- comments begin with `#` and are thus ignored
- multi line comments can be created with `#=` and terminated with `#=`
- colored output can be created with `print_with_color(:red, "I love Julia!")`
- objects are often interacted with in Julia – actions on objects are written functionally, like `action(object)` and not `object.action()`
- to display objects from code while outside of REPL, use `display(object)`

Types

- the type system is pretty unique, variables can be bound again to the same name
- if a variable is given a type, it will only accept variables of the same type
- type annotations are generally done as `var::TypeName` – type annotations make error checking and optimization easier

```
calculate_position(time::Float64)
```

- type assertions work in exactly the same way – an error is raised if the types differ

```
(expression)::TypeName
```

- to achieve top performance, types of variables have to remain stable
- Julia has a type hierarchy, with `Any` being the top – any undeclared variable is of this type
- the type `Nothing` is one no object can have, but it is a subtype of all other types
- custom types can be defined – their named in CamelCase

- `typeof(x)` gives the type of variables, `isa(x, T)` checks if `x` is of type `T`
- data types are all of type `DataType`
- if data types can be converted by using the lower case type name

```
int64(3.14)
```

Integers and Booleans

- available integers are `Int8`, `Int16`, `Int32`, `Int64`, `Int128`
- unsigned integers are defined using `U` in front of the type name, like `UInt32`
- depending on machine architecture, the standard is either 32 or 64 bit
- this is given by the variable `WORD_SIZE`
- `typemin(var)` and `typemax(var)` give the min and max sizes of types
- if `typemax` is exceeded, overflow occurs
- explicit checking for overflow needs to occur
- other ways of writing integers are: `0b` (binary), `0o` (octal), `0x` (hexadecimal)
- in case arbitrary sizes are needed, `BigInt` can be used

```
BigInt("1234")
```

- they support the same operations as normal integers, but conversions are not automatic
- divisions of `BigInt` always give floating point results
- `div` gives integer division, `rem` gives the remainder, `^` gives powers
- `true` and `false` are actually 8bit integers, where `0` is `false` and anything else is `true`, `bool(32)` returns `true`
- other logical operations are `!`, `==`, `!=`, `<`, `>`
- logical expressions can be chained `0 < x < 3`

Floating point numbers

- follow IEEE 754
- come in `Float16`, `Float32`, `Float64`
- single precision floating point numbers: `2.5f2 == 2.5 * 10^2`
- double precision floating point numbers: `2.5e2 == 2.5 * 10^2`
- `BigFloat` is used for arbitrary precision floats
- `Inf` is infinity, `NaN` is not a number
- floating point numbers do cause errors because of the bad precision

Elementary mathematical functions and operations

- the bit representation of any number can be seen using the `bits(num)` operation
- `round()` rounds to float
- `iround()` rounds to int
- other functions: `sqrt()`, `cbrt()`, `exp()`, `log()`, `sin()`, `cos()`, `tan()`, `erf()` (error function), `rand()`
- `()` enforce precedence
- chained assignments are allowed `a = b = c = d = 1`
- assignments can be combined `a, b = c, d` means `a = c` and `b = d`

- bool operations have `||` and `&&`, `|` and `&` are used for non-short-circuit operations, those that check all the options even if the result is already found
- bitwise operations are supported <http://docs.julialang.org/en/latest/manual/mathematical-operations/>

Rational and complex numbers

- i is represented as `im`, `3.2 + 7.1im` has the type `Complex{Float64}`
- all normal maths functions are also implemented for this type
- `abs()` returns the absolute value
- `complex(a, b)` turns `a` and `b` into a complex number
- rational numbers are good for exact ratios
- `3//4` is a rational for $\frac{3}{4}$ and has the type `Rational{Int64}`
- `float()` converts to a floating point number, `num()` and `denum()` give the numerator and the denominator

Characters

- `'a'` is a char with type `Char` – a 32-bit integer with a unicode code point inside of it
- `int('A')` returns 65, `char(65)` returns `A`
- unicode can be entered using single quotes around `'\uxxxx'` or `'\Uxxxxxxxx'`
- `is_valid_char()` checks if the entered number is a valid character
- normal escape sequences also exist, `\n`, `\t`,...

Strings

- strings are always ASCII or UTF8, depending on the contained letters
- strings are immutable and contained in `"string"` or `'''string'''`
- arrays of chars starting at 1, so `s[1]` returns the first letter
- `s[end]` returns the end, `endof(s)` returns the last byte's index, `length()` the length
- iteration through a string is best done in an iteration and not an index

```
for c in string
    println(c)
end
```

- substrings can be obtained by string slices, `s[3:5]` or `s[2:end]`
- string interpolation

```
a = 3; b = 2;
s = "$a * $b = $(a * b)"
# s is "3 * 2 = 6"
```

- concatenate with

```
"abd" * "def" == string("abc", "def")
```

- **Symbols** are strings prefixed with a colon `:green`
- these more efficient strings are used as IDs or keys, they can't be concatenated
- `methodswith(String)` returns all the functions that use it

- `search(string, char)` – returns the index of the first char in the string
- `replace(string, str1, str2)` – replaces `str1` in `string` with `str2`
- `split(string, char or [chars])` – splits `string` at `char`, returns an array of strings
– if there is no `char` given, `split` uses whitespace

Formatting numbers and strings

- formatted printing is done with the `@printf` macro, it takes one format string and one or more variables

```
name = "Paskal"
@printf("Hello, %s\n", name)
# prints "Hello, Paskal"
```

- to return the formatted string, use `@sprintf`
- `show()` prints a text representation of an object that is often more specific than simple `@printf`
 - integer `%d`

```
@printf("%d\n", 100)           #> 100
```

- float `%.3f`

```
@printf("%.2f\n", 2.346)       #> 2.34
# or
str = @sprintf("%.2f\n", 2.346)
show(str)                      #> "2.34"
```

- scientific `%0.3e`

```
@printf("%.2e", 123)           #> 1.23e+2
```

- character `%c`

```
@printf("%c", 'd')             #> d
```

- string `%s`

```
@printf("%s", "hi")            #> hi
```

- right justify `%50s`

```
@printf("%8s", "hi")           #>          hi
```

- a special type of string is the `VersionNumber`, preceded by a `v`, is used to run different code for different versions of Julia

```
v"0.3"
```

Regular Expressions

- Julia uses Perl syntax for regex, see <http://www.regular-expressions.info/reference.html>
- regular expressions are denoted by a string preceded by `r`

```
r"..."
```

- matching an email address

```
# '+' matches any non-empty group of characters
email_pattern = r".+@.+"
input = "john.doe@mit.edu"
println(ismatch(email_pattern, input)) #> true
```

- `ismatch(pattern, string)` returns a boolean whether or not the pattern is matched
- `match(patterns, string)` returns nothing or `RegexMatch` if the pattern is found
- `RegexMatch` has the properties `match` (contains substring that matches), `offset` (position of match in string), `offsets` (start positions for each of the substrings), `captures` (captured substrings as tuple)
- capturing in regex can be done with parenthesis `()`, line

```
email_pattern = r"(.+)@(.+)"
```

- this will capture the hostname and username of the email address
- `replace` also takes regex as inputs to replace by
- `matchall` returns all the matches in an array
- `eachmatch` returns an iterator of matches

Ranges and arrays

- ranges work as

```
for i in <start>:(<interval>):<end>
    ...
end
````
- the start and end are inclusive, if the interval is one it can be omitted
- arrays are of type Array and can only hold the same type of value
- `Array{T, n}` is type and dimension
- 1D arrays are also called vectors
````
julia
arr = [100, 25, 37]
```

- arrays are really vectors and they grow dynamically
- arrays can also be constructed abstractly

```
arr = Array{Int64, 5}
show(arr) #> [0,0,0,0,0]
```

- to add to an array, use `push!`

```
push!(arr, 1)
```

- arrays can also be created from slices

```
arr = [1:9]
```

- to allocate a certain amount of memory (so it does not need to be resized) to one array at once do

```
sizehint(arr, 10^5)
```

- array indices start at 1 and can be used in brackets `arr[1]` and `end` is the final element
- `eltype` gives element type, `length` the number of elements, `ndims` the number of dimensions, `size(arr, n)` gives the elements per dimension
- to turn an array into a comma separated string, use `join(arr, "sep")`
- slices can retrieve sub-arrays and they can also be assigned to

```
arr[1:3]    #> [1, 2, 3]
arr[1:3] = 1
arr[1:3] = [2, 3, 4]
```

Other ways to create arrays

- `zeros(n)` returns an array with `n` 0s, `ones` does the same for 1s
- `linspace(start, stop, n)` returns an array of `n` with equally spaced numbers from `start` to `stop`
- `fill!(arr, n)` returns an array filled with only `ns`
- to get random values, use

```
arr = rand{T, n}
# returns n random numbers of type T
```

- arrays can also be converted using the normal conversion functions

Some common functions for arrays

- concatenate arrays with `append!(arr1, arr2)`
- **any function ending in ! modifies the first element it is given**
- `push!` and `pop!` work like on the stack
- `shift!` pops the first element, `unshift!` pushes an element to `arr[1]`
- `splice!(index)` removes the element at `index`
- `in(val, arr)` checks if `val` is in `arr`
- `sort` and `sort!` sort the arrays, the latter changes it
- for loops can iterate over arrays
- element wise operations are done with a point

```
a1 = [1, 2, 3]
a1 = [2, 4, 6]
a1 .* a2
```

- the dot product is done with `dot(a1, a2)`
- the sum over an array is `sum(a1, a2)`
- arrays are passed as references – they point to the same memory
- many functions are simply applied to each element in an array
- `repeat` repeats each element in an array a certain number of times
- `copy` makes a copy of the first level, `deepcopy` goes through everything recursively

Array of chars to string

- `join(arr)` turns a `[char]` into a string, `utf32(arr)` does the same

- to pass the contents of array as separate parameters, use the splice operator, like `string(arr...)`

Dates and times

- `time()` returns seconds since epoch
- `strftime(time())` returns more useful formatted time
- `Date` and `DateTime` functions return more useful and subtractable values

Scope and constants